

Moore_Week11_12

Heather

2025-02-20

Introduction to Machine Learning

Importing the data

```
binary = read.csv("C:/Users/14027/Documents/Graduate_Schoolish/DSC_520/Code_Hmwk/binary-classifier-data.csv")
trinary = read.csv("C:/Users/14027/Documents/Graduate_Schoolish/DSC_520/Code_Hmwk/trinary-classifier-data.csv")
```


Plot the data from each data set using a scatter plot

```
library(GGally)


## Loading required package: ggplot2

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2

library(ggplot2)
#Correlation of binary
#bin = ggpairs(binary)
ggplot(binary, aes(x,y)) + geom_point()
```



```
#Correlation og trinary  
ggplot(trinary, aes(x,y)) +geom_point()
```



In this problem, you will determine which points are nearest by calculating the Euclidean distance between two points. As a refresher, the Euclidean distance between two points

```
library(stats)
#binary
vect1 = binary$x
vect2 = binary$y
CalculateEuclideanDistance <- function(vect1, vect2) sqrt(sum((vect1 - vect2)^2))

print("Euclidean distance between vect1 and vect2 from the binary dataset is: ")
```

```
## [1] "Euclidean distance between vect1 and vect2 from the binary dataset is: "
```

```
# Calling CalculateEuclideanDistance function
CalculateEuclideanDistance(vect1, vect2)
```

```
## [1] 1411.959
```

```
#trinary
vect1 = trinary$x
vect2 = trinary$y
```

```
CalculateEuclideanDistance <- function(vect1, vect2) sqrt(sum((vect1 - vect2)^2))

print("Euclidean distance between vect1 and vect2 from the trinary dataset is: ")
```

```
## [1] "Euclidean distance between vect1 and vect2 from the trinary dataset is: "
```

```
# Calling CalculateEuclideanDistance function
CalculateEuclideanDistance(vect1, vect2)
```

```
## [1] 1357.734
```

Fit a k nearest neighbors' model for each dataset for k=3, k=5, k=10, k=15, k=20, and k=25. Compute the accuracy of the resulting models for each value of k. Plot the results in a graph where the x-axis is the different values of k and the y-axis is the accuracy of the model

```
#Split into test and train sets in each data set
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(e1071)
```

```
## Warning: package 'e1071' was built under R version 4.2.2
```

```
library(class)
train_index = sample(nrow(binary), size = floor(nrow(binary) * 0.75))
## Binary
bin_train = binary[train_index, ]
bin_test = binary[-train_index,]

## Trinary
tri_train = trinary[train_index,]
tri_test = trinary[-train_index,]
```

Fit the Binary Model

```

#For loop for Binary
k_vals = c(3,5,10,15,20,25)
bin_accuracy = data.frame(k = k_vals, accuracy = NA)

# Loop through each k value
for (i in 1:length(k_vals)) {
  k = k_vals[i]

  # Perform k-NN classification
  bin = knn(train = bin_train, test = bin_test, cl = bin_train$label, k = k)

  # Get true labels
  labels = bin_test$label

  # Calculate accuracy
  bin_acc = mean(labels == bin)

  # Store the accuracy in the data frame
  bin_accuracy$accuracy[i] = bin_acc
}

# Print the final data frame
print(bin_accuracy)

```


```

##      k  accuracy
## 1   3 0.9733333
## 2   5 0.9680000
## 3  10 0.9706667
## 4  15 0.9706667
## 5  20 0.9733333
## 6  25 0.9733333

```

Model Graph

```
plot(bin_accuracy)
```



Creation of the Trinary Model

```
#For loop for Trinary
k_vals = c(3,5,10,15,20,25)
tri_accuracy = data.frame(k = k_vals, accuracy = NA)
# Loop through each k value
for (i in 1:length(k_vals)) {
  k = k_vals[i]

  # Perform k-NN classification
  bin = knn(train = tri_train, test = tri_test, cl = tri_train$label, k = k)

  # Get true labels
  labels = tri_test$label

  # Calculate accuracy
  bin_acc = mean(labels == bin)


  # Store the accuracy in the data frame
  tri_accuracy$accuracy[i] = bin_acc
}

# Print the final data frame
print(tri_accuracy)
```

```
##      k  accuracy
## 1   3 0.8022472
## 2   5 0.7932584
## 3  10 0.7842697
## 4  15 0.7685393
## 5  20 0.7550562
## 6  25 0.7617978
```

Creating the Plot

```
plot(tri_accuracy)
```



Do you think a linear classifier would work well on these data sets?

- Based on the the scatter plot from above it shows that there is no obvious linear relationship in the data in either data set. So I would say that the linear classifiers would not work well.


**How does accuracy of your logistic regression classifier from last week compare?
Why is the accuracy different between these two methods?**

The two classifiers last week had an accuracy of 54% and 83% which are smaller than what the accuracies are from this week. The logistic regression was supposed to predict a linear model but, the data isn't linear

while the KNN predicts the closest data points in the training set which worked a lot better for both data set this time around.

Clustering


```
library(stats)
# Importing the Data
cluster = read.csv("C:/Users/14027/Documents/Graduate_Schoolish/DSC_520/Code_Hmwk/clustering-data.csv")
#Plot the data set in a scatter plot
ggplot(cluster, aes(x,y)) +geom_point()
```





Fit the dataset using the k-means algorithm from k=2 to k=12. Create a scatter plot of the resultant clusters for each value of k.


```
k_group = c(2:12)
k_dist = data.frame()
for(x in k_group) {
  set.seed(123)
  cluster_model = kmeans_results = kmeans(cluster, x, nstart = 10)
  cluster$cluster = as.factor(cluster_model$cluster)
  plot = ggplot(cluster, aes(x,y, color = cluster)) + geom_point()


  print(plot)
}
```




























Calculate this average distance from the center of each cluster for each values of k and plot it as a line chart where k is the x-axis and the average distance is the y-axis.

```
# Loop through each k value
set.seed(123) # Set seed for reproducibility

# Run k-means clustering
cluster_model = kmeans(cluster, centers = 12, nstart = 10)

# Store the centers (convert to list to store in a data frame)
k_center = list(cluster_model$centers)
k_center = as.data.frame(k_center)
k_center$avg_distance = sqrt((k_center$y - k_center$x)^2)

ggplot(k_center, aes(x = cluster, y = avg_distance)) +
  geom_line()
```



One way of determining the “right” number of clusters is to look at the graph of k versus average distance and finding the “elbow point”. Looking at the graph you generated in the previous example, what is the elbow point for this dataset?

- My graph does not follow the proper curve of a elbow curve. However after doing some research the elbow point is where k has reached the point of diminishing returns. The closest point that matches that definition in my opinion would be $k=8$.